

Cost of the prior vs. accuracy of the estimate (live LLM pool)
generation 5.36 s/item/model measured on gpt-5.6-luna; MAE is the mean over 3 modalities

